

Here's a full draft of documentation content designed for both Python and npm SDKs, formatted for Mintlify (using MDX/Markdown).

---

## Introduction

```
---
```

```
title: "Introduction"
description: "Welcome to Agnitra AI — runtime optimization for AI models"
---
```

```
# Welcome to **Agnitra AI**
```

Agnitra AI is the agentic runtime optimisation engine for your AI models. Whether you're running large language models, diffusion, or other tensor-heavy workloads — Agnitra helps you:

- Run your model **\*\*faster\*\*** (higher tokens/sec, lower latency)
- Use **\*\*less memory\*\*** and compute
- Support **\*\*any hardware\*\*** (NVIDIA, AMD, Tenstorrent, etc.)
- Do it with **\*\*no code rewrites\*\*** — just integrate via SDK or CLI

> “One-line integration. Runtime patching. Cross-chip.”

```
---
```

## ## Why this matters

Modern AI models are massive and expensive to operate. Many users face:

- Manual tuning of kernels & performance bottlenecks
- Vendor-locked optimisation stacks
- Hard to scale or port across different accelerators
- Slow iteration cycles when trying to tune performance

Agnitra addresses these by embedding runtime telemetry, LLM-/RL-based optimisation, and kernel patching under the hood — letting you focus on your model, not the infrastructure.

```
---
```

## ## What this docs cover

This documentation covers:

- Installing and integrating via **\*\*Python SDK\*\***, **\*\*npm/TS SDK\*\***, or **\*\*CLI\*\***
- Using our REST API if you choose the SaaS or hosted path
- Walk-throughs for use-cases (LLMs, diffusion, edge)
- Backend architecture overview
- Enterprise deployment & hardware targets
- FAQ and troubleshooting

Let's get started.

---

## Getting Started

```
---  
title: "Getting Started"  
sidebarTitle: "Getting Started"  
description: "Install and run your first model optimisation in minutes."  
---
```

### ## Quickstart

#### ### Prerequisites

- Node.js version 16+ for the npm path
- Python 3.8+ for the Python SDK
- A model file (e.g., `llama.pt`, `model.onnx`) or telemetry log
- Access to your target hardware (GPU/accelerator)

#### ### Install

##### #### Python

```
``bash  
pip install agnitra
```

##### npm / TypeScript

```
npm install agnitra  
# or  
yarn add agnitra
```

---

## First optimisation

### Python SDK

```
from agnitra import optimize_model  
model = load_model("llama.pt") # load your model however you prefer  
optimized = optimize_model(model, target="A100")
```

### npm / TypeScript SDK

```
import { optimizeModel } from "agnitra";  
const model = await loadModel("llama.pt");  
const optimized = await optimizeModel(model, { target: "A100" });
```

 That's it — after optimisation you can run optimized in place of your original model for inference.

---

### Using the CLI

```
agnitra optimize --model llama.pt --target A100
```

This will profile, optimise, generate kernels, and patch at runtime.

---

## Next steps

- Explore the SDK Reference for full API usage
- Check the CLI Reference for advanced options
- Review the Use-Cases section to see how others use Agnitra

---

```
# SDK Reference
```

```
```mdx
```

```
---
```

```
title: "SDK Reference"
sidebarTitle: "SDK Reference"
description: "Detailed API for Python and npm SDKs."
```

```
---
```

```
# Python SDK
```

```
## `optimize_model(model, *, target: str, mode?: str) -> model`  
Optimise a loaded model for the given hardware target.
```

**\*\*Parameters\*\***

- `model`: your loaded model (PyTorch, ONNX, etc.)
- `target`: e.g., `"A100"`, `"H100"`, `"MI250"`,
- `mode` (optional): `"throughput" | "memory"` etc.

**\*\*Returns\*\***: Optimised model instance.

```
## `benchmark_model(original_model, optimized_model) -> Dict[str, float]`  
Compare before/after metrics: tokens/sec, latency, memory.
```

```
---
```

```
# npm / TypeScript SDK
```

```
## `optimizeModel(model: Model, options: { target: string; mode?: string }) : Promise<Model>`  
Same functionality as Python, adapted for JS/TS.
```

```
```ts
```

```
import { optimizeModel } from "agnitra";
```

```
const optimized = await optimizeModel(myModel, { target: "A100", mode: "throughput" });
```

```
benchmarkModel(original: Model, optimized: Model) : Promise<BenchmarkResult>
```

Returns object such as { tokensPerSecGain: 0.28, latencyReduction: 0.21, memoryReduction: 0.18 }.

---

CLI Reference

(Insert detailed CLI flags, sub-commands, examples)

---

```
# Use Cases
``mdx
---
title: "Use Cases"
sidebarTitle: "Use-Cases"
description: "Real-world model optimisation examples"
---
```

## LLaMA-7B on A100

We profiled LLaMA-7B with baseline tokens/sec = 320. After optimisation we achieved ~410 (+28%) on A100, latency dropped ~20%, memory ~18%.

#### Python SDK example

```
``python
from agnitra import optimize_model
opt = optimize_model(model, target="A100", mode="throughput")
```

npm SDK example

```
import { optimizeModel } from "agnitra";
const opt = await optimizeModel(model, { target: "A100", mode: "throughput" });
```

Stable Diffusion XL (SDXL) on H100

...

---

```
# Architecture
``mdx
---
title: "Architecture"
sidebarTitle: "Architecture"
description: "How Agnitra works under the hood"
---
```

## Pipeline Overview

1. Telemetry Collector – captures GPU time, memory, tensor shapes
2. IR Graph Extraction – builds intermediate representation (e.g., torch.fx)
3. LLM + RL Agent – suggests kernel strategies; learns over time
4. Kernel Generator & Runtime Patcher – builds Triton/CUDA kernels, injects at runtime
5. Feedback Loop & Benchmarking – captures results, refines intelligence

## Hardware Support

List of supported chips, versioning details...

## Data Flow Diagram

(Insert architecture diagram here)

---

## FAQ

---

title: "FAQ"  
sidebarTitle: "FAQ"  
description: "Common questions and answers"

---

### What kinds of models are supported?

We support PyTorch, ONNX, Tensor Flow in most cases. Telemetry ingestion is flexible.

### What hardware is supported?

Any GPU/accelerator that has accessible profiling and kernel generation capabilities (NVIDIA-A100/H100, AMD-MI series, Tensorflow etc.)

### Is the optimisation safe?

Yes — we generate kernels tested in sandbox; your original model remains unchanged until you choose to swap in the optimised version.

### What about enterprise/on-prem?

See our \*Enterprise Deployment\* section for VPC, private GPU clusters, compliance and security.

---

## Enterprise Deployment

---

title: "Enterprise Deployment"  
sidebarTitle: "Enterprise"  
description: "On-prem, VPC, OEM licensing"

---

## Deployment Models

- **Hosted SaaS** — easy to adopt
- **On-Prem SDK** — installs in your environment
- **OEM Embedding** — integrate with hardware vendors

## Compliance & Security

Details on data privacy, telemetry anonymisation, and certs.

## Custom Domain & Domain-Path Setup

We support publishing docs at `docs.yourcompany.com` or `/docs` path.

## Support & SLA

Enterprise customers receive dedicated SLAs, support, custom model tuning, annual training etc.

---

## CLI Reference

```
---
```

title: "CLI Reference"  
sidebarTitle: "CLI"  
description: "Command-line usage and flags"  
---

## `agnitra optimize`

agnitra optimize --model --target [-mode throughput|memory] [-telemetry ]

### Flags

- `--model` : path to your model file
- `--target` : hardware target (e.g., A100)
- `--mode` : optimisation mode
- `--telemetry` : (optional) supply pre-collected telemetry JSON

### Examples

```
``bash
agnitra optimize --model llama.pt --target A100
agnitra optimize --model sdxl.onnx --target H100 --mode memory
```

---

agnitra benchmark

agnitra benchmark --before model.pt --after optimized\_model.pt

Generates a report, e.g., tokens/sec, latency, memory.

```
---
```

### Configuration File — `docs.json`  
This file configures your Mintlify docs site layout, themes, nav structure.

```
```json
{
  "name": "Agnitra AI Docs",
  "logo": {
    "light": "/logo/light.svg",
    "dark": "/logo/dark.svg"
  },
  "colors": {
    "primary": "#00C4A7",
    "background": "#111827"
  },
  "sidebar": [
    { "title": "Getting Started", "items": ["introduction", "getting-started"] },
    { "title": "SDK Reference", "items": ["sdk-reference"] },
    { "title": "CLI Reference", "items": ["cli-reference"] },
    { "title": "Use-Cases", "items": ["use-cases"] },
    { "title": "Architecture", "items": ["architecture"] },
  ]
}
```

```
{ "title": "Enterprise", "items": ["enterprise-deployment"] },
{ "title": "FAQ", "items": ["faq"] }
],
"editUrl": "https://github.com/yourorg/docs/edit/main/docs/",
"favicon": "/favicon.ico"
}
```

---